

PyTorch

And Some Things About It

Viet Nhat Nguyen - 21520378

Summer Training
University of Information Technology

Jul. 12, 2023

Summary

1 What is PyTorch?

2 Tensors

3 Autograd

4 Basics

Warm-up: numpy

Warm-up: Numpy

We can implement the network using numpy, but

- it does not know anything about computation graphs, or deep learning, or gradients.
- start from scratch, or manually!

Warm-up: Numpy

We can implement the network using numpy, but

- it does not know anything about computation graphs, or deep learning, or gradients.
- start from scratch, or manually!

Warm-up: Numpy

We can implement the network using numpy, but

- it does not know anything about computation graphs, or deep learning, or gradients.
- start from scratch, or manually!

PyTorch

PyTorch is a python package that provides two high-level features:

- An n-dimensional Tensor, similar to numpy
- Automatic differentiation for building and training neural networks

PyTorch

PyTorch is a python package that provides two high-level features:

- An n-dimensional Tensor, similar to numpy but **can run on GPUs**
- Automatic differentiation for building and training neural networks

Why PyTorch?

- More Pythonic (imperative)
 - Flexible
 - Intuitive and cleaner code
 - Easy to debug
- More Neural Networkic
 - Write code as the network works
 - forward/backward

Install PyTorch

PyTorch Build	Stable (2.0.1)		Preview (Nightly)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 11.7	CUDA 11.8	ROCm 5.4.2	CPU
Run this Command:	<pre>pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu117</pre>			

Creating Tensors

Creating Tensors

So many ways to create a Tensor:

- `torch.empty()`
- `torch.zeros()`
- `torch.ones()`
- `torch.rand()`
- `torch.tensor()`

```
import torch  
my_tensor = torch.tensor([2, 10, 23])
```

Tensor Data Types

Tensor Data Types

dtype argument:

- `torch.bool`
- `torch.int8`
- `torch.uint8`
- `torch.int16`
- `torch.int32`
- `torch.int64`
- `torch.half`
- `torch.float`
- `torch.double`
- `torch.bfloat`

Tensor Shapes

Tensor Shapes

```
my_tensor.shape
```


Tensor Shapes

Sometimes, tensors need to be of the same *shape* - that is, having the same number of dimensions and the same number of cells in each dimensions.

- We have `torch.*_like()` methods

Tensor Shapes

Sometimes, tensors need to be of the same *shape* - that is, having the same number of dimensions and the same number of cells in each dimensions.

- We have `torch.*_like()` methods

Math & Logic with PyTorch Tensors

Math & Logic with PyTorch Tensors

- Basic arithmetic: $+$ $-$ $*$ $/$ \dots
 - Tensor with Scalar
 - Tensor with Tensor

As we can see, it is element-wise

Math & Logic with PyTorch Tensors

- Basic arithmetic: $+$ $-$ $*$ $/$ \dots
 - Tensor with Scalar
 - Tensor with Tensor

As we can see, it is element-wise

Math & Logic with PyTorch Tensors

- An **element-wise** operation operates on corresponding elements between tensors.

$$\begin{bmatrix} a_{1,1} & a_{1,2} \\ a_{2,1} & a_{2,2} \\ a_{3,1} & a_{2,2} \end{bmatrix} \diamond \begin{bmatrix} b_{1,1} & b_{1,2} \\ b_{2,1} & b_{2,2} \\ b_{3,1} & b_{2,2} \end{bmatrix} = \begin{bmatrix} a_{1,1} \diamond b_{1,1} & a_{1,2} \diamond b_{1,2} \\ a_{2,1} \diamond b_{2,1} & a_{2,2} \diamond b_{2,2} \\ a_{3,1} \diamond b_{3,1} & a_{3,2} \diamond b_{3,2} \end{bmatrix}$$

- But, what happens when we try to perform operations on tensors with dissimilar shape?

Math & Logic with PyTorch Tensors

- An **element-wise** operation operates on corresponding elements between tensors.

$$\begin{bmatrix} a_{1,1} & a_{1,2} \\ a_{2,1} & a_{2,2} \\ a_{3,1} & a_{2,2} \end{bmatrix} \diamond \begin{bmatrix} b_{1,1} & b_{1,2} \\ b_{2,1} & b_{2,2} \\ b_{3,1} & b_{2,2} \end{bmatrix} = \begin{bmatrix} a_{1,1} \diamond b_{1,1} & a_{1,2} \diamond b_{1,2} \\ a_{2,1} \diamond b_{2,1} & a_{2,2} \diamond b_{2,2} \\ a_{3,1} \diamond b_{3,1} & a_{3,2} \diamond b_{3,2} \end{bmatrix}$$

- But, what happens when we try to perform operations on tensors with dissimilar shape?

Math & Logic with PyTorch Tensors

- Basic arithmetic: $+$ $-$ $*$ $/$ \dots
 - *Tensor with Scalar*
 - *Tensor with Tensor (same-shape)*
 - **Tensor with Tensor (not same-shape)***

Math & Logic with PyTorch Tensors

The exception to the same-shape rule is *tensor broadcasting*

Math & Logic with PyTorch Tensors

- **Broadcasting** is a way to perform an operation between tensors that have similarities in their shapes.
 - multiplying a tensor of learning weights by a *batch* of input tensors
- Broadcasting's rules:
 - Each tensor must have at least one dimension - no empty tensors.
 - Comparing the dimension sizes of the two tensors, *backward*:
 - Each dimension must be equal, *or*
 - One of the dimensions must be of size 1, *or*
 - The dimension does not exist in one of the tensor.
 - Tensors of the identical shape.

Math & Logic with PyTorch Tensors

- **Broadcasting** is a way to perform an operation between tensors that have similarities in their shapes.
 - multiplying a tensor of learning weights by a *batch* of input tensors
- Broadcasting's rules:
 - Each tensor must have at least one dimension - no empty tensors.
 - Comparing the dimension sizes of the two tensors, *backward*:
 - Each dimension must be equal, *or*
 - One of the dimensions must be of size 1, *or*
 - The dimension does not exist in one of the tensor.
 - Tensors of the identical shape.

Math & Logic with PyTorch Tensors

- **Broadcasting** is a way to perform an operation between tensors that have similarities in their shapes.
 - multiplying a tensor of learning weights by a *batch* of input tensors
- Broadcasting's rules:
 - Each tensor must have at least one dimension - no empty tensors.
 - Comparing the dimension sizes of the two tensors, *backward*:
 - Each dimension must be equal, *or*
 - One of the dimensions must be of size 1, *or*
 - The dimension does not exist in one of the tensor.
 - Tensors of the identical shape.

Math & Logic with PyTorch Tensors

- **Broadcasting** is a way to perform an operation between tensors that have similarities in their shapes.
 - multiplying a tensor of learning weights by a *batch* of input tensors
- Broadcasting's rules:
 - Each tensor must have at least one dimension - no empty tensors.
 - Comparing the dimension sizes of the two tensors, *backward*:
 - Each dimension must be equal, *or*
 - One of the dimensions must be of size 1, *or*
 - The dimension does not exist in one of the tensor.
 - Tensors of the identical shape.

Math & Logic with PyTorch Tensors

- **Broadcasting** is a way to perform an operation between tensors that have similarities in their shapes.
 - multiplying a tensor of learning weights by a *batch* of input tensors
- Broadcasting's rules:
 - Each tensor must have at least one dimension - no empty tensors.
 - Comparing the dimension sizes of the two tensors, *backward*:
 - Each dimension must be equal, *or*
 - One of the dimensions must be of size 1, *or*
 - The dimension does not exist in one of the tensor.
 - Tensors of the identical shape.

Math & Logic with PyTorch Tensors

- **Broadcasting** is a way to perform an operation between tensors that have similarities in their shapes.
 - multiplying a tensor of learning weights by a *batch* of input tensors
- Broadcasting's rules:
 - Each tensor must have at least one dimension - no empty tensors.
 - Comparing the dimension sizes of the two tensors, *backward*:
 - Each dimension must be equal, *or*
 - One of the dimensions must be of size 1, *or*
 - The dimension does not exist in one of the tensor.
 - Tensors of the identical shape.

Math & Logic with PyTorch Tensors

- *Basic arithmetic: $+$ $-$ $*$ $/$...*
 - *Tensor with Scalar*
 - *Tensor with Tensor (same-shape)*
 - *Tensor with Tensor (not same-shape)**
- More Math with Tensors

Moving to GPU

Autograd

Automatic Differentiation

Automatic Differentiation




Figure: Training a RNN

Automatic Differentiation

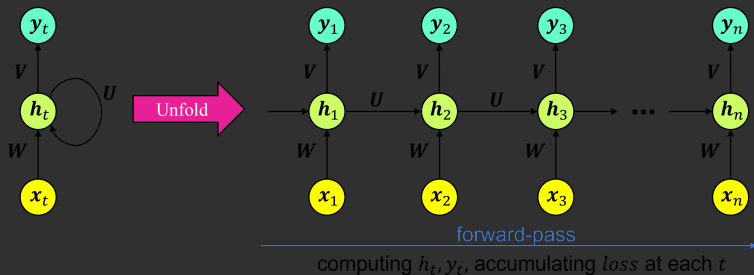


Figure: Training a RNN

Automatic Differentiation

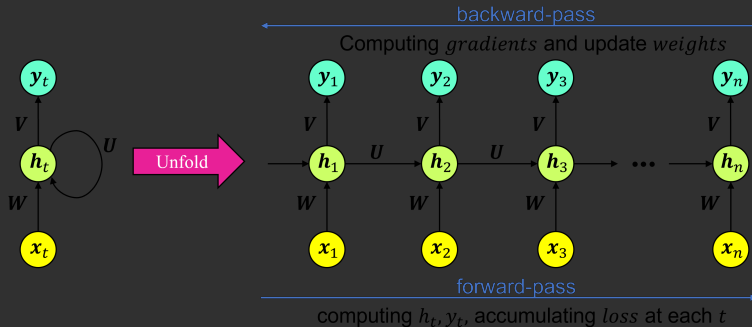


Figure: Training a RNN

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 -

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 -

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 - $\frac{\partial L}{\partial x} = \frac{\partial L(y)}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial M(x)}{\partial x}$

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 - $\frac{\partial L}{\partial x} = \frac{\partial L(y)}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial M(x)}{\partial x}$

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 - $\frac{\partial L}{\partial x} = \frac{\partial L(y)}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial M(x)}{\partial x}$

Automatic Differentiation

- Machine learning model: $y = M(x)$
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 - $\frac{\partial L}{\partial x} = \frac{\partial L(y)}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial M(x)}{\partial x}$

Automatic Differentiation

- Machine learning model: $y = M(x)$.
- The *loss* function: $L(y) = L(M(x))$
 - Our goal is to *minimize* the loss
 - Nudge the weights until a tolerable loss for wide variety of inputs
 - $$\frac{\partial L}{\partial x} = \frac{\partial L(y)}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial L}{\partial y} \frac{\partial M(x)}{\partial x}$$

Simple Autograd

Practice

Vector Calculus using autograd

Vector Calculus using autograd

$$\vec{y} = f(\vec{x})$$

The Jacobian matrix J :

$$J = \left(\frac{\partial y}{\partial x_1} \cdots \frac{\partial y}{\partial x_n} \right) = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_n} \end{pmatrix}$$

The gradient of a scalar function $l = g(\vec{y})$:

$$v = \left(\frac{\partial l}{\partial y_1} \cdots \frac{\partial l}{\partial y_m} \right)^T$$

Vector Calculus using autograd

$$\vec{y} = f(\vec{x})$$

The Jacobian matrix J :

$$J = \left(\frac{\partial y}{\partial x_1} \cdots \frac{\partial y}{\partial x_n} \right) = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_n} \end{pmatrix}$$

The gradient of a scalar function $l = g(\vec{y})$:

$$v = \left(\frac{\partial l}{\partial y_1} \cdots \frac{\partial l}{\partial y_m} \right)^T$$

Vector Calculus using autograd

$$\vec{y} = f(\vec{x})$$

The Jacobian matrix J :

$$J = \left(\frac{\partial y}{\partial x_1} \cdots \frac{\partial y}{\partial x_n} \right) = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_n} \end{pmatrix}$$

The gradient of a scalar function $l = g(\vec{y})$:

$$v = \left(\frac{\partial l}{\partial y_1} \cdots \frac{\partial l}{\partial y_m} \right)^T$$

Vector Calculus using autograd


By the chain rule, gradient of l w.r.t \vec{x} :

$$J^T \cdot v = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_n} & \cdots & \frac{\partial y_m}{\partial x_n} \end{pmatrix} \begin{pmatrix} \frac{\partial l}{\partial y_1} \\ \vdots \\ \frac{\partial l}{\partial y_m} \end{pmatrix} = \begin{pmatrix} \frac{\partial l}{\partial x_1} \\ \vdots \\ \frac{\partial l}{\partial x_n} \end{pmatrix}$$

Note: We can also use the equivalent operation $v^T \cdot J$


Computational Graphs

Computational Graphs




Chain rule: $\frac{\partial f}{\partial a} = \frac{\partial f}{\partial d} \times \frac{\partial d}{\partial a} = e \times 1 = e$

Computational Graphs



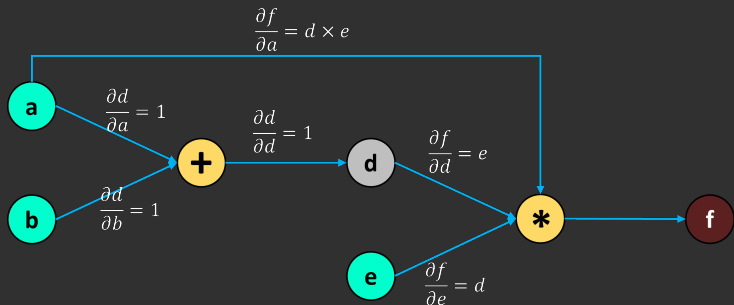
Chain rule: $\frac{\partial f}{\partial a} = \frac{\partial f}{\partial d} \times \frac{\partial d}{\partial a} = e \times 1 = e$

Computational Graphs



Chain rule:
$$\frac{\partial f}{\partial a} = \frac{\partial f}{\partial d} \times \frac{\partial d}{\partial a} = e \times 1 = e$$

Computational Graphs



Chain rule:
$$\frac{\partial f}{\partial a} = \frac{\partial f}{\partial d} \times \frac{\partial d}{\partial a} + \frac{\partial f}{\partial a} = e \times 1 + d \times e = (e + 1) \times d$$

Computational Graphs

In the forward pass:

- run the requested operation to compute a resulting tensor, and
- maintain the operation's *gradient function* in the DAG.

The backward pass kicks off when `.backward()`:

- computes the gradients from each `.grad_fn`,
- accumulates them in the respective tensor's `.grad` attribute, and
- using the chain rule, propagates all the way to the leaf tensors.

Building a Backwards Graph

Hooks

Other Basics

Dataset and DataLoaders

Save and Load the Model

The End